

Homework 2

solutions by Luigi Russo, 1699981

1 Hashing

(a) Let $\mathcal{H} = \{H_s : \{0,1\}^{2n} \rightarrow \{0,1\}^n\}_{s \in \{0,1\}^\lambda}$ be a family of collision-resistant hash functions compressing $2n$ bits into n bits. Answer the following questions.

(i) Show that \mathcal{H} is a seeded one-way function in the following sense: for all PPT adversaries \mathbf{A} there exists a negligible function $\nu : \mathbb{N} \rightarrow [0, 1]$ such that

$$\mathbb{P}[H_s(x') = y : s \leftarrow \{0,1\}^\lambda; x \leftarrow \{0,1\}^{2n}; y \leftarrow H_s(x); x' \leftarrow \mathbf{A}(s, y)] \leq \nu(n)$$

SOL

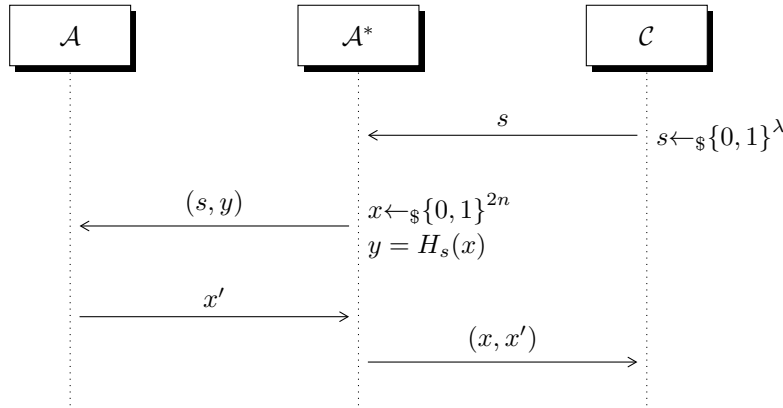


Figure 1: Base an attack to CR \mathcal{H} family, assuming the existence of a PPT attacker \mathcal{A} , able to win the seeded one-way function Game with non negligible probability.

Assume there exists a PPT attacker \mathcal{A} that given a value $y = H_s(x)$, with $s \leftarrow_{\$} \{0,1\}^\lambda$ and $x \leftarrow_{\$} \{0,1\}^{2n}$, finds a pre-image $x' : H_s(x) = H_s(x')$ with non negligible probability $\geq \frac{1}{p(n)}$ for some $p(n) \in \text{poly}(n)$. Then we can build on top of \mathcal{A} an attacker \mathcal{A}^* able to break the collision resistance of \mathcal{H} .

In particular, the challenger selects and fixes a seed s ; \mathcal{A}^* fixes a value x and forwards to \mathcal{A} the value $y = H_s(x)$; \mathcal{A} returns a value x' and the tuple (x, x') is forwarded to the challenger: the game is won if $x \neq x'$ and $H_s(x) = H_s(x')$. Now, by construction, we know that \mathcal{A} "inverts"¹ y with probability $\geq \frac{1}{p(n)}$. We also know that the domain of \mathcal{H} is the set of $2n$ -bit strings, while its co-domain is the set of n -bit strings: this implies that the expected value of elements $x : H_s(x) = y$, once we fix y and s , is 2^n , thus implying

¹with some abuse, we mean that it finds a valid preimage of y .

that the probability for the attacker \mathcal{A} to output a value x' equal to x is only negligible ($= 2^{-n}$). These two considerations allow to conclude that \mathcal{A}^* outputs a valid and colliding pair (x, x') with probability $\geq \frac{1}{p(n)}(1 - 2^{-n}) \in \text{poly}(n)$. \square

- (ii) What happens in case the set of functions \mathcal{H} is not compressing (i.e., the domain of each function H_s is also $\{0, 1\}^n$? Does collision resistance imply one-wayness in this case?

SOL Consider $H_s(x) := x \oplus \text{pad}(s)$, where $\text{pad} : \{0, 1\}^\lambda \rightarrow \{0, 1\}^n$ is public padding function². Then it is easy to see that \mathcal{H} is collision resistant, since there is no collision at all (for each s , $H_s(\cdot)$ defines a permutation). By the way, this family of functions is definitely not one way; it is easily invertible, since the seed s is public in the game, i.e. it is known to the adversary too: indeed, given $y = H_s(x) = \text{pad}(s) \oplus x$, the attacker outputs $\mathcal{A}(s, y) = y \oplus \text{pad}(s) = \text{pad}(s) \oplus \text{pad}(s) \oplus x = x$. \square

- (b) Let $\mathcal{H} = \{H_s : \{0, 1\}^{4n} \rightarrow \{0, 1\}^{2n}\}_{s \in \{0, 1\}^\lambda}$ and $\mathcal{H}' = \{H'_s : \{0, 1\}^{2n} \rightarrow \{0, 1\}^n\}_{s \in \{0, 1\}^\lambda}$ be families of collision-resistant hash functions. Analyse the following candidate hash function family compressing $4n$ bits into n bits: $\mathcal{H}^* := \{H_{s_1, s_2}^* : \{0, 1\}^{4n} \rightarrow \{0, 1\}^n\}_{s_1, s_2 \in \{0, 1\}^\lambda}$ such that $H_{s_1, s_2}^*(x) = H'_{s_2}(H_{s_1}(x))$ for $s_1, s_2 \leftarrow_{\$} \{0, 1\}^\lambda$.

SOL

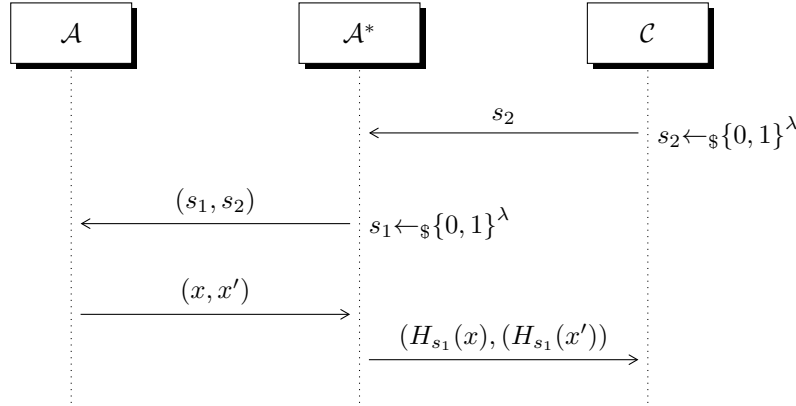


Figure 2: Base an attack to CR \mathcal{H}' family on top of an attacker \mathcal{A} to the collision resistance of \mathcal{H}^* .

Let assume that \mathcal{H}^* is not collision resistant: this means that, by definition, there exists some PPT attacker \mathcal{A} that, given two values as seed s_1, s_2 , is able to find, with non negligible probability, a valid pair (x, x') such that $H_{s_1, s_2}^*(x) = H_{s_1, s_2}^*(x')$, with $x \neq x', x, x' \in \{0, 1\}^{4n}$. If this is the case we can try to build an attacker \mathcal{A}^* to break the collision resistance of \mathcal{H}' . Note that the challenger fixes a seed s_2 and waits for the attacker to find a collision; \mathcal{A}^* randomly selects and fixes a seed $s_1 \in \{0, 1\}^\lambda$ and uses \mathcal{A} in order to find a pair (x, x') . By assumption, this pair is such that $H_{s_1, s_2}^*(x) = H_{s_1, s_2}^*(x')$ with non negligible probability $\geq \frac{1}{p(n)}$, for some $p(n) \in \text{poly}(n)$.

But if \mathcal{A} is successful (i.e. $H_{s_1, s_2}^*(x) = H_{s_1, s_2}^*(x')$), either (i) $H_{s_1}(x) = H_{s_1}(x')$ or (ii) $H_{s_1}(x) \neq H_{s_1}(x')$. Case (i) can happen only with negligible probability $\epsilon(n)$, otherwise we could use \mathcal{A} to build a PPT attacker to break the collision resistance of \mathcal{H} . Instead case (ii) guarantees that the pair is valid to build a successful attack to \mathcal{H}' : in both situations we come to an absurd given our initial assumptions. \square

²e.g. $\text{pad}(s) := s || 0^{n-\lambda}$.

2 Number Theory

- (a) Recall that the CDH problem asks to compute g^{ab} given (g, g^a, g^b) for $(G, g, q) \leftarrow \text{\texttt{GroupGen}}(1^\lambda)$ and $a, b \leftarrow \mathbb{Z}_q$. Prove that the CDH problem is equivalent to the following problem: given (g, g^a) compute g^{a^2} , where $(G, g, q) \leftarrow \text{\texttt{GroupGen}}(1^\lambda)$ and $a \leftarrow \mathbb{Z}_q$.

SOL For sake of simplicity, let us call SCDH the problem defined above (i.e. compute g^{a^2} , given a tuple (\mathbb{G}, g, q, g^a)). We have to prove that both $\text{CDH} \Rightarrow \text{SCDH}$ and $\text{SCDH} \Rightarrow \text{CDH}$.

(CDH \Rightarrow SCDH) given some g^a it is easy to see that this is equivalent to solve $\frac{\text{CDH}(g, g^a, g^{a+r})}{g^{ar}} = \frac{g^{aa+ar}}{g^{ar}} = g^{a^2}$, for a random choiche of $r \leftarrow \mathbb{Z}_q$.

(SCDH \Rightarrow CDH) given (g^a, g^b) , we can compute $\frac{\text{SCDH}(g^{a+b})}{\text{SCDH}(g^a) \cdot \text{SCDH}(g^b)} = \frac{g^{a^2+2ab+b^2}}{g^{a^2} \cdot g^{b^2}} = g^{2ab}$: at this point we need to compute the square root of such a value: generally speaking this operation cannot be done efficiently; however, there are some groups where computing the square root can be done very easily (e.g. cyclic groups or when the order q is odd). For example, consider the case for q odd: it holds that $(g^{2ab})^{\frac{q+1}{2}} = 1 \cdot (g^{2ab})^{\frac{1}{2}} = g^{ab}$.

□

- (b) Let $f_{g,p} : \mathbb{Z}_{p-1} \rightarrow \mathbb{Z}_p^*$ be the function defined by $f_{g,p}(x) := g^x \bmod p$. Under what assumption is $f_{g,p}$ one-way? Prove that the predicate $h(x)$ that returns the least significant bit of x is not hard-core for $f_{g,p}$.

SOL The function $f_{g,p}(\cdot)$ is one-way if g is a generator of \mathbb{Z}_p^* , p is prime and under the DL assumption. First of all we know that this function is efficiently computable: this is a requirement for all OWF. On the contrary it is hard to "invert", i.e. to find a value $x' : f(x) = f(x')$. Indeed, given $y = g^x \bmod p$, a PPT attacker should find a value $x' : g^x \bmod p = g^{x'} \bmod p$. But this is possible if and only if $x - x' \equiv 0 \bmod (p-1)$.

Note that an element $g^x \in \mathbb{QR}_p$ if and only if the least significant bit of the exponent x is 0. Without loss of generality, indeed, we assume $x \in \{0, 1\}^t$: this means that x can be written as $\sum_{i=0}^t x_i \cdot 2^i$. Given $f_{g,p}(x) = g^x \bmod p = g^{\sum_{i=0}^t x_i \cdot 2^i} \bmod p$, we can check whether $f_{g,p}(x) \in \mathbb{QR}_p$ or not, in order to leak the least significant bit of x .

The last step is to find a necessary and sufficient condition to check whether an element y is a quadratic residue modulo p or not.

We can compute $s = y^{\frac{p-1}{2}}$; if $s \equiv 1 \bmod p$, then y is a quadratic residue; otherwise it is not.

(\Rightarrow) If y is a quadratic residue, then it holds that $y = (g^\alpha)^2 = g^{2\alpha}$, for some α . But then: $y^{\frac{p-1}{2}} = g^{\alpha(p-1)} = (g^{p-1})^\alpha = 1^\alpha = 1 \bmod p$.

(\Leftarrow) If it holds that $y^{\frac{p-1}{2}} \equiv 1 \bmod p$, then $g^{\frac{x(p-1)}{2}} \equiv 1 \bmod p$. This implies that $\frac{x(p-1)}{2} \equiv 0 \bmod (p-1)$ that can be satisfied if and only if x is even. Finally, we conclude that $y = g^x = (g^{\frac{x}{2}})^2$, proving that y is a quadratic residue.

□

- (c) Let N be the product of 5 distinct odd primes. If $y \in \mathbb{Z}_N^*$ is a quadratic residue, how many solutions are there to the equation $x^2 = y \bmod N$?

SOL $x^2 \equiv y \pmod{N}$ means that $x^2 \equiv y \pmod{p_1 \cdot p_2 \cdot p_3 \cdot p_4 \cdot p_5}$. This implies that there exists a $k \in \mathbb{Z}$ such that $y = x^2 + k \prod_{i=1}^5 p_i$. If we define $y_i := y \pmod{p_i}$, we can rewrite the equation as $y_i \equiv x^2 \pmod{p_i}, i \in \{1, \dots, 5\}$. This system of equation can be safely rewritten as $x^2 \equiv y_i \pmod{p_i}$, but now, for the Chinese remainder theorem, we know that once we fix the values y_i for all i , there exists a unique solution satisfying the system³. Note that each y_i is such that it is equivalent (modulo p_i) to x_i^2 , where $x_i := x \pmod{p_i}$. Each of the 5 equations has two distinct solutions $\pm \alpha_i$, since $\alpha_i \not\equiv -\alpha_i \pmod{p_i}$, because $p_i \neq 2, \forall i$. And this proves that there are 2^5 distinct systems, each of these with a unique solution (for CR theorem). This implies that the number of solutions to the original equation is $2^5 = 32$. \square

3 Public-Key Encryption

- (a) Show that for any CPA-secure public-key encryption scheme for single-bit messages, the length of the ciphertext must be super-logarithmic in the security parameter.

SOL If the length of the ciphertext is not super-logarithmic, we know that the number of possible ciphertexts is at most λ^c for some constant c . The attacker can compute a ciphertext c_0 associated to bit 0, then waits for the challenge c_b^* and compares the two values: if they are equal it returns 0, otherwise returns 1. Now, if the challenger selects $b = 0$, the attacker wins with probability at least λ^{-c} , otherwise wins with probability 1. Taking the weighted average we have a winning probability $w \geq \frac{1}{2} + \frac{1}{2\lambda^c}$.

But then $|w - \frac{1}{2}| \geq \frac{1}{2\lambda^c}$, where $2\lambda^c \in \text{poly}(\lambda)$. \square

- (b) Let $\Pi = (\text{KGen}, \text{Enc}, \text{Dec})$ be a PKE scheme with message space $\{0, 1\}$ (i.e., for encrypting a single bit). Consider the following natural construction of a multi-bit PKE scheme $\Pi' = (\text{KGen}', \text{Enc}', \text{Dec}')$ with message space $\{0, 1\}^t$, for some polynomial $t = t(\lambda)$: (i) The key generation stays the same, i.e. $\text{KGen}'(1^\lambda) = \text{KGen}(1^\lambda)$; (ii) Upon input $m = (m[1], \dots, m[t]) \in \{0, 1\}^t$ the encryption algorithm $\text{Enc}'(pk, m)$ outputs a ciphertext $c = (c_1, \dots, c_t)$ where $c_i \leftarrow_{\$} \text{Enc}(pk, m[i])$ for all $i \in [t]$; (iii) Upon input a ciphertext $c = (c_1, \dots, c_t)$ the decryption algorithm $\text{Dec}'(sk, c)$ outputs the same as $(\text{Dec}(sk, c_1), \dots, \text{Dec}(sk, c_t))$.

- (i) Show that if Π is CCA1 secure, so is Π' .

SOL In order to prove the above claim, it is convenient to use the hybrid argument: note that the goal is to prove that for all PPT adversaries \mathcal{A} , it holds that $\text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 0) \approx_c \text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 1)$.

In particular let us define a set of intermediate hybrid games Hyb_i , where the adversary is given access to the decryption oracle (together with all the public parameters to make encryptions too, of course) exactly as each of the two games above: the difference is that, when the adversary commits the challenge tuple (m_0, m_1) , the challenge ciphertext c^* is such that its j -th component is sampled from $\text{Enc}(pk, m_0[j])$ if $i \leq j$, otherwise comes from $\leftarrow_{\$} \text{Enc}(pk, m_1[j])$.

Note that once we fix the pair (m_0, m_1) , the only difference between two hybrids Hyb_i and Hyb_{i+1} is in the distribution of the i -th component of the challenge c^* . It is easy to see that $\text{Hyb}_1 \equiv \text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 0)$ and $\text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 1) \equiv \text{Hyb}_t$.

Here for simplicity we give the distribution of the challenge ciphertext c^* for each hybrid game:

- $\text{Hyb}_1 \rightarrow (\text{Enc}(pk, m_0[1]), \text{Enc}(pk, m_0[2]), \dots, \text{Enc}(pk, m_0[t])) \equiv \text{Enc}'(m_0)$

³with a simple substitution, we put $T := x^2$ and can apply the theorem.

- $\text{Hyb}_2 \rightarrow (\text{Enc}(pk, m_1[1]), \text{Enc}(pk, m_0[2]), \dots, \text{Enc}(pk, m_0[t]))$
- \dots
- $\text{Hyb}_{t-1} \rightarrow (\text{Enc}(pk, m_1[1]), \dots, \text{Enc}(pk, m_1[t-1]), \text{Enc}(pk, m_0[t]))$
- $\text{Hyb}_t \rightarrow (\text{Enc}(pk, m_1[1]), \dots, \text{Enc}(pk, m_1[t-1]), \text{Enc}(pk, m_1[t])) \equiv \text{Enc}'(m_1)$

Imagine there exists a PPT distinguisher between two hybrids Hyb_i and Hyb_{i+1} (for some i). If this is the case, we can build a PPT attacker \mathcal{A} to break the CCA1-security of Π . Indeed, we can ask the challenger to give us an encryption c_b^* of a bit b . Then we pass to the distinguisher a value c such that the first i components are $\leftarrow_{\$} \text{Enc}(pk, m_0[j]), \forall j < i$, its i -th component is equal to c_b^* and then the last $t-i-1$ components are $\leftarrow_{\$} \text{Enc}(pk, m_1[j])$. We then forward to the challenger the same bit returned by the distinguisher and we retain the same non negligible advantage⁴. This proves that $\forall i, \text{Hyb}_i \approx_c \text{Hyb}_{i+1}$ which implies that also $\text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 0) \approx_c \text{Game}_{\Pi, \mathcal{A}}^{\text{CCA-1}}(\lambda, 1)$. \square

- (ii) Show that, even if Π is CCA2 secure, Π' is not CCA2 secure.

SOL Even if the original scheme is CCA2, it is very simple to break the CCA2-security of Π' : the attacker obtains a ciphertext c_b^* where the first bit of m_0^* is 0 as well as the first bit of m_1^* . Now it produces another ciphertext for the bit 0: $x := \text{Enc}(0)$ such that x is different from the encryption of the first bit of m_b^* , and then asks for the decryption of $c' = (x || c_b^{*5})$, and finally returns the decision bit $b' : m_{b'}^* = \text{Dec}(c')$, winning with probability 1. \square

- (c) Consider the following variant of El Gamal encryption. Let $p = 2q + 1$, let \mathbb{G} be the group of squares modulo p (so \mathbb{G} is a subgroup of \mathbb{Z}_p^* of order q), and let g be a generator of \mathbb{G} . The private key is (\mathbb{G}, g, q, x) and the public key is (\mathbb{G}, g, q, h) , where $h = g^x$ and $x \in \mathbb{Z}_q$ is chosen uniformly. To encrypt a message $m \in \mathbb{Z}_q$, choose a uniform $r \in \mathbb{Z}_q$, compute $c_1 := g^r \bmod p$ and $c_2 := h^r + m \bmod p$, and let the ciphertext be (c_1, c_2) . Is this scheme CPA-secure? Prove your answer.

SOL This scheme is not CPA-secure. Indeed, note that in order to compute c_1 we add an element $h^r \in \mathbb{G}$ to the message $m \in \mathbb{Z}_q$. This means that $c_2 \in \mathbb{Z}_p^*$ but not necessarily it is an element of \mathbb{G} : in particular, an attacker could exploit the fact that choosing uniformly a message $m \leftarrow_{\$} \mathbb{Z}_q$, $\mathbb{P}[c_2 \in \mathbb{G} : c_2 = h^r + m \bmod p; h^r \in \mathbb{G}] = \frac{q}{2q+1}$, because among all the $2q+1$ elements of the group, only q of them are squares (and then belong to \mathbb{G}).

Note also that when $m = 0$, $c_2 = h^r + m = h^r \Rightarrow c_2 \in \mathbb{G}$.

So our attacker can use these two messages $m_0^* = 0$ and $m_1^* \leftarrow_{\$} \mathbb{Z}_q$ and can receive the challenge $c_b^*(c_1^*, c_2^*)$: if $c_2^* \in \mathbb{G}$ it outputs 0, otherwise outputs 1. Note that if $b = 1$, the attacker wins with probability 1; otherwise it wins with probability $\frac{q+1}{2q+1}$. But this is enough to retain a non negligible advantage (the winning probability is $w = \frac{1}{2} + \frac{q+1}{2(2q+1)} \Rightarrow |w - \frac{1}{2}|$ is non negligible in λ). \square

4 Signature Schemes

- (a) Consider a weaker variant of UF-CMA in which the attacker receives (pk, m^*) at the beginning of the experiment, where the message m^* is uniformly random over $\mathbb{Z}_{\mathbb{N}}^*$, and thus it has to

⁴this proof is not valid if $m_0[i] = m_1[i]$: but in such a scenario a distinguisher between Hyb_i and Hyb_{i+1} cannot even exist, since the i -th component is sampled from the same distribution and then the two distributions are equivalent.

⁵it is c_b^* without the first block, the one corresponding to the encryption of the first bit of m_b^*

forge on m^* after possibly seeing polynomially-many signatures σ_i on uniformly random messages $m_i \leftarrow \$_{\mathbb{Z}_N^*}$ chosen by the challenger. Call this notion random-message unforgeability under random-message attacks (RUF-RMA).

Formalize the above security notion, and prove that UF-CMA implies RUF-RMA but not viceversa.

SOL Given a signature scheme $\Pi := (\text{KGen}, \text{Sign}, \text{Vrfy})$, let consider the following game $\text{Game}_{\Pi, \mathcal{A}}^{\text{RUF-RMA}}(\lambda)$:

- (a) the challenger \mathcal{C} generates $(pk, sk) \leftarrow \$_{\text{KGen}}(1^\lambda)$
- (b) the challenger \mathcal{C} fixes a message $m^* \leftarrow \$_{\mathbb{Z}_N^*}$ and forwards to \mathcal{A} both the public key pk and the message m^*
- (c) \mathcal{A} has access to oracle $O_1(sk)$: this oracle samples a message $m_i \leftarrow \$_{\mathbb{Z}_N^*}$ and returns (m_i, σ_i) such that $\text{Vrfy}(pk, (m_i, \sigma_i)) = 1$
- (d) \mathcal{A} forges the signature σ^*
- (e) the output of the game is $\text{Vrfy}(pk, (m^*, \sigma^*))$ (i.e. 1 in case of win, 0 otherwise)

We say that a scheme Π is RUF-RMA-secure if \forall PPT attackers \mathcal{A} , $\exists \epsilon(\lambda) \in \text{negl}(\lambda)$:

$$\mathbb{P}[\text{Game}_{\Pi, \mathcal{A}}^{\text{RUF-RMA}}(\lambda) = 1] \leq \epsilon(\lambda)$$

(UF-CMA \Rightarrow RUF-RMA) Assume that a scheme $\Pi := (\text{KGen}, \text{Sign}, \text{Vrfy})$ is not RUF-RMA: then there exists a PPT attacker \mathcal{A} able to forge, with non negligible probability, a valid signature σ^* for the message m^* fixed by the challenger. If this is true, we can build on top of \mathcal{A} an attacker \mathcal{A}^* able to forge a valid pair and break UF-CMA-security too.

First of all, in this reduction, \mathcal{A}^* forwards to \mathcal{A} all the public parameters. Moreover, it randomly selects a message m^* and relays this to \mathcal{A} : m^* simulates the random challenge. At this point, the attacker \mathcal{A} can start querying the oracle O_1 ; but \mathcal{A}^* can simulate very easily these queries, simply picking a message at random $m_i \leftarrow \$_{\mathbb{Z}_N^*}$ and asking the challenger \mathcal{C} to forge a valid signature σ_i : the pair (m_i, σ_i) is then forwarded to \mathcal{A} . After potentially a polynomial number of these oracle calls, \mathcal{A} is ready to forge a signature σ^* and \mathcal{A}^* forwards to \mathcal{C} the pair (m^*, σ^*) , retaining the same non-negligible advantage.

(RUF-RMA $\not\Rightarrow$ UF-CMA) The textbook version of RSA signatures is RUF-RMA (see next exercise), however it is not UF-CMA as proved in class. \square

- (b) Recall the textbook-version of RSA signatures.

$\text{KGen}(1^\lambda)$: Run $(N, e, d) \leftarrow \$_{\text{GenModulus}}(1^\lambda)$, and let $pk = (e, N)$ and $sk = (N, d)$.

$\text{Sign}(sk, m)$: Output $\sigma = m^d \bmod N$.

$\text{Vrfy}(pk, m, \sigma)$: Output 1 if and only if $\sigma^e \equiv m \bmod N$.

Prove that the above signature scheme $\Pi = (\text{KGen}, \text{Sign}, \text{Vrfy})$ satisfies RUF-RMA under the RSA assumption.

SOL Assume not: then there exists a PPT attacker \mathcal{A} (against RUF-RMA of Π), on top of which we can break RSA assumption. In particular, the challenger generates RSA params (N, e, d) and fixes, at random, a value $x \leftarrow \$_{\mathbb{Z}_N^*}$ and computes $y = x^e$. The attacker \mathcal{A}^* forwards the public key (N, e) to \mathcal{A} ; moreover it fixes $m^* = y$. In order to simulate the calls to oracle O_1 it can do the following: generates a signature $\sigma_i \leftarrow \$_{\mathbb{Z}_N^*}$ and then computes the message $m_i = \sigma_i^e$ and returns the pair (m_i, σ_i) . Note that such a pair is valid (because $\sigma_i^e = m_i$ by definition,

and since σ_i is chosen at random, also m_i looks random⁶). We know that \mathcal{A} forges a valid pair (y, σ^*) with non negligible probability, such that $\sigma^* : \sigma^{*e} = y \Rightarrow \sigma^* = x$. This concludes the proof. \square

5 Actively Secure ID Schemes

Let $\Pi = (\text{KGen}, \text{P}, \text{V})$ be an ID scheme. Informally, an ID scheme is actively secure if no efficient adversary \mathcal{A} (given just the public key pk) can make V accept, even after \mathcal{A} participates maliciously in poly-many interactions with P (where the prover is given both the public key pk and the secret key sk). More formally, we say that Π satisfies active security if for all PPT adversaries \mathcal{A} there is a negligible function $\nu : \mathbb{N} \in \{0, 1\}$ such that:

$$\mathbb{P}[\text{Game}_{\Pi, \mathcal{A}}^{\text{mal-id}}(\lambda) = 1] \leq \nu(\lambda),$$

where the game $\text{Game}_{\Pi, \mathcal{A}}^{\text{mal-id}}(\lambda)$ is defined as follows:

- The challenger runs $(pk, sk) \leftarrow_{\$} \text{KGen}(1^\lambda)$, and returns pk to \mathcal{A} .
- Let $q(\lambda) \in \text{poly}(\lambda)$ be a polynomial. For each $i \in [q]$, the adversary can run the protocol Π with the challenger (where the challenger plays the prover and the adversary plays the malicious verifier), obtaining transcripts $\tau_i \leftarrow_{\$} (\text{P}(pk, sk) \rightleftharpoons \mathcal{A}(pk))$.
- Finally, the adversary tries to impersonate the prover in an execution of the protocol with the challenger (where now the challenger plays the honest verifier), yielding a transcript $\tau^* \leftarrow_{\$} (\mathcal{A}(pk) \rightleftharpoons \text{V}(pk))$.
- The game outputs 1 if and only if the transcript τ^* is accepting, i.e. $\text{V}(pk, \tau^*) = 1$.

Answer the following questions.

- (a) Prove that passive security is strictly weaker than active security. Namely, show that every ID scheme Π that is actively secure is also passively secure, whereas there exists a (possibly contrived) ID scheme Π_{bad} that is passively secure but not actively secure.

SOL (active \Rightarrow passive) Assume not, then there exists some PPT attacker \mathcal{A} that breaks the passive security of an actively secure scheme Π . If this is true, we can build a new PPT attacker \mathcal{A}^* that breaks the active security of Π . Indeed, when \mathcal{A} asks to see some valid transcript τ_i , prompting the empty string, \mathcal{A}^* will simply start a session of the protocol as an **honest**⁷ verifier with the challenger acting as the (honest) prover. The final transcript τ_i is then forwarded to the original attacker: and this can be done up to a polynomial number of times of course. At this point, \mathcal{A} impersonates a malicious prover (because it has not the secret key sk), directly interacting with the verifier (\mathcal{A}^* , of course, forwards the messages of \mathcal{A} to the challenger and then sends its replies back to \mathcal{A}). This allows \mathcal{A}^* to retain the same non negligible advantage of \mathcal{A} since the reduction is tight.

(passive $\not\Rightarrow$ active) Consider a canonical Σ protocol Π : we define a variant Π' such that the verifier appends a bit b to its (unique) message β where $\mathbb{P}[b = 0] = 2^{-\lambda}$; if $b = 0$ (this happens only with negligible probability in λ) the prover has to send the secret key sk ; otherwise it follows

⁶this is crucial, otherwise it would not be a valid simulation.

⁷must behave honestly, for each transcript requested by \mathcal{A} , in order to perfectly simulate the oracle.

the original protocol to compute γ ⁸. This clearly does not compromise the passive security of Π' , because the attacker should hope to see a transcript where the secret key is leaked, but this happens only with negligible probability since the verifier is honest (it is simulated by the challenger itself); on the contrary, this scheme is not actively secure, because the adversary is allowed to impersonate the verifier in the first part of its attack: then it can easily find out the secret key (needs only one query where it appends $b = 0$ to some β) and then it can act as a malicious prover and forge a valid transcript thanks to the leaked key. \square

- (b) Let $\Pi' = (\text{KGen}, \text{Sign}, \text{Vrfy})$ be a signature scheme, with message space \mathcal{M} . Prove that if Π' is UF-CMA, the following ID scheme $\Pi = (\text{KGen}, \text{P}, \text{V})$ (based on Π') achieves active security:

$\text{P}(pk, sk) \Rightarrow \text{V}(pk)$: The verifier picks random $m \leftarrow_{\$} \mathcal{M}$, and forwards m to the prover. The prover replies with $\sigma \leftarrow_{\$} \text{Sign}(sk, m)$, and finally the verifier accepts if and only if $\text{Vrfy}(pk, m, \sigma) = 1$.

SOL As always, assume not: then if there exists a PPT \mathcal{A} able to break the active security of Π , we can build on top of it a new attacker \mathcal{A}^* to break the UF-CMA-security of Π' . In our reduction, we first forward to the original attacker the public key pk (and, of course, all the public parameters generated by \mathcal{C}).

The attacker \mathcal{A} is allowed in the first phase to interact as a (possibly malicious) verifier: it will start sending some message m_i , waiting for the prover to reply; \mathcal{A}^* will forward m_i to the challenger in order to produce a valid σ_i and then is ready to reply; this allows the attacker \mathcal{A} to collect some transcript τ_i . After polynomially many such queries, it tries to impersonate the prover: so \mathcal{A}^* will pick a random (and "fresh"⁹, i.e. $\neq m_i, \forall i$) message m^* and passes it to \mathcal{A} ; then, a σ^* is forged (valid with probability $\geq \frac{1}{p(\lambda)}, p(\lambda) \in \text{poly}(\lambda)$); the pair (m^*, σ^*) is forwarded to the challenger and the winning condition is that σ^* is a valid signature of m^* . But this directly implies that \mathcal{A}^* has the same non negligible advantage of the original \mathcal{A} . \square

- (c) Is the above protocol honest-verifier zero-knowledge? Prove your answer.

SOL In order to be HVZK, it should exist a simulator S such that $\{pk, sk, \text{Trans}(pk, sk)\} \approx_c \{pk, sk, S(pk)\}$. Note that this simulator is given in input only the public key pk . But if such S exists, the underlying signature scheme Π cannot be UF-CMA: this because S is able to produce valid pairs (m_i, σ_i) in a way that is (computationally) indistinguishable from $\text{Trans}(pk, sk)$ and in particular it can forge a valid pair in $\text{Game}_{\Pi'}^{\text{UF-CMA}}(\lambda)$ too.

If, instead, we remove the UF-CMA assumption of Π' , it could be that for some schemes it is possible to compute pairs (m_i, σ_i) given only pk : e.g. a possible strategy could be to pick a random σ_i and then find some message m_i such that $\text{Vrfy}(pk, (m_i, \sigma_i)) = 1$. If this is the case, then we can build such a simulator S and the above protocol is HVZK. \square

⁸of course V' will check the last bit of β and if it is equal to 1, it will run $\text{V}(\cdot)$.

⁹note that it can do it, since it observes all the queries m_i .